

Core Buildings Blocks of RAG:


Core Building Blocks of RAG

RAG systems are powered by essential components like embeddings, vector stores, retrievers, chunking strategies, and prompt engineering. Understanding these building blocks is key to designing scalable, accurate, and domain-aware AI assistants.

01

Core Building Blocks of RAG

- Embeddings
- Vector Stores
- Text Splitters & Chunking
- Retrievers
- Prompt Engineering for RAG



User Query → [Embed] → [Retrieve] → [Inject into Prompt] → [LLM Generates Answer]

↑ ↑
Vector DB Prompt Template

02