

Vector Stores

Vector stores are specialized databases that store embeddings and support fast similarity search. They are critical in RAG pipelines for retrieving the most relevant context chunks based on a user's query.

01

Vector Stores

What it is

- Databases that store **vector embeddings** and support **similarity search**.

Why it's important

- Acts as the knowledge base that stores indexed documents in vector format, enabling efficient search.

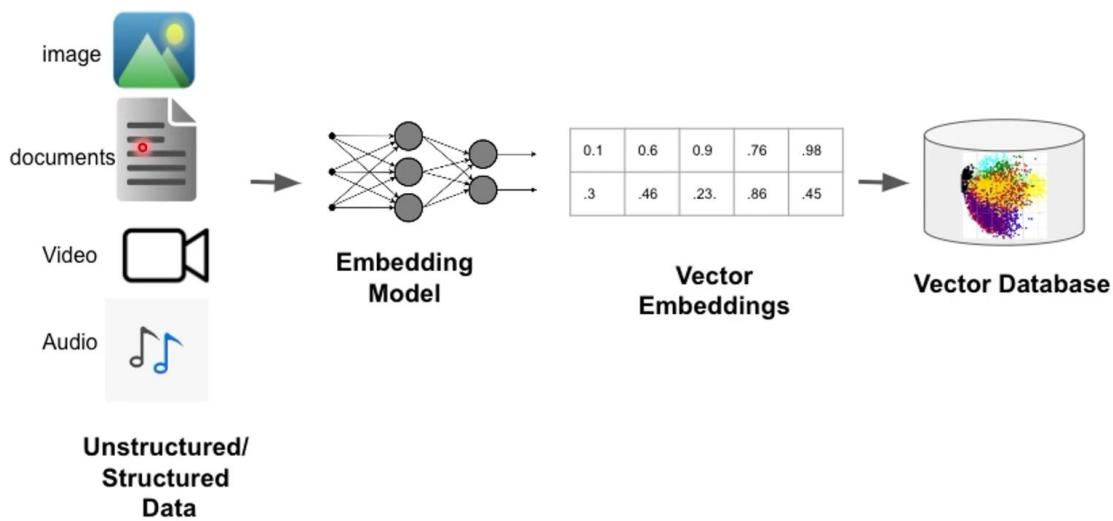
Examples

- FAISS (Facebook AI Similarity Search)
- Pinecone
- Chroma
- Weaviate
- Qdrant



02

What is Vector Database



03

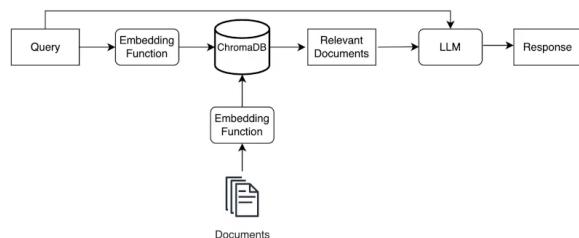
What is ChromaDB?

ChromaDB – Lightweight Local Vector Store for AI Apps

- ❑ Open-source vector database written in Python.
- ❑ Ideal for local development and prototyping.
- ❑ Supports document storage with metadata.
- ❑ Integrated with LangChain and LlamaIndex.
- ❑ Backed by DuckDB for persistence (optional).



Chroma



04

Key Features of ChromaDB

- ❑ **Simple API:** Add, update, and query vectors with ease.
- ❑ **Local-first:** Runs in memory or with disk-based DuckDB.
- ❑ **Document-aware:** Supports documents, metadata, and embeddings.
- ❑ **Python-native:** No setup or external services needed.
- ❑ **Integrations:** Works seamlessly with LangChain and LlamaIndex.



05

Use Cases for ChromaDB

- ❑ Local RAG pipelines (e.g., PDF/chatbots).
- ❑ Streamlit or notebook-based assistants.
- ❑ AI research tools and semantic memory stores.
- ❑ MVPs and offline demos.



06