

Pinecone DB:

What is Pinecone?

Pinecone – Cloud-Native Vector Database for Scalable Search

- ❑ Fully managed vector database as a service.
- ❑ Designed for real-time, low-latency semantic search.
- ❑ Supports billions of vectors with namespaces and metadata.
- ❑ Optional hybrid search: dense + sparse indexing.



07

Key Features of Pinecone

- ❑ **Serverless & Scalable:** No infra to manage.
- ❑ **Real-time search:** <100ms latency at scale.
- ❑ **Namespaces & Metadata:** Great for multi-tenant apps.
- ❑ **Hybrid Search:** Combines semantic and keyword-based results.
- ❑ **Integrated with OpenAI, Hugging Face, LangChain.**



08

Use Cases for Pinecone

- ❑ Production RAG systems with OpenAI/Anthropic.
- ❑ Semantic search in SaaS and web apps.
- ❑ Multi-user AI chatbots with memory.
- ❑ Search over enterprise knowledge bases.

